

MTIL Arguments Guide

This document provides a comprehensive reference for all command-line arguments available in the MTIL (Multi-Task Incremental Learning) training pipeline.

Table of Contents

- [Testing](#)
 - [Hyperparameters](#)
 - [Logging & Training Control](#)
 - [Experiment Settings](#)
 - [Single Image Evaluation](#)
 - [Save & Load](#)
 - [Model Control \(image-fc branch\)](#)
 - [ZSCL \(Zero-Shot Continual Learning\)](#)
 - [iCaRL](#)
 - [Loss Functions](#)
 - [WiSE-FT \(Weight-Space Ensembling\)](#)
 - [Experiment Organization](#)
 - [Model Freezing](#)
 - [Learning Without Forgetting \(LwF\)](#)
 - [Baseline & TRIO](#)
 - [Fisher Weighting](#)
 - [Thesis-Specific Arguments](#)
 - [LoRA \(Low-Rank Adaptation\)](#)
-

Testing

Argument	Type	Default	Description
--test	flag	False	Enable testing mode

Hyperparameters

Argument	Type	Default	Description
--model	str	ViT-B/16	CLIP model architecture. Options: ViT-B/16, ViT-B/32, ViT-L/14, RN50, RN101
--batch-size	int	8	Training batch size
--batch-size-eval	int	32	Evaluation batch size
--lr	float	0.001	Learning rate
--wd	float	0.0	Weight decay for optimizer

Argument	Type	Default	Description
--ls	float	0.0	Label smoothing factor
--warmup_length	int	100	Number of warmup iterations for learning rate scheduler
--beta2	float	0.999	Beta2 parameter for Adam optimizer

Logging & Training Control

Argument	Type	Default	Description
--seed	int	42	Random seed for reproducibility
--epochs	int	None	Number of training epochs (mutually exclusive with --iterations)
--iterations	int	None	Number of training iterations (mutually exclusive with --epochs)
--eval-interval	int	None	Evaluate every N iterations (mutually exclusive with --eval-every-epoch)
--loss-interval	int	1000	Log loss every N iterations
--eval-every-epoch	flag	False	Evaluate at the end of each epoch
--eval-only	flag	False	Skip training, only run evaluation
--save-eval	flag	False	Save evaluation results
--start-iteration	int	None	Starting iteration (for resuming training)

Experiment Settings

Argument	Type	Default	Description
--method	str	finetune	Training method. Choices: finetune, lwf, ZSCL, icarl
--train-mode	str	whole	Which parts of the model to train. Choices: whole, text, image, image-fc, image-fc-fixed, fc
--data-location	str	./data	Root directory for datasets
--train-dataset	str	None	Dataset to train on (e.g., DTD, CIFAR100, ImageNet)

Argument	Type	Default	Description
--eval-datasets	str	None	Comma-separated list of datasets for evaluation
--text-datasets	str	None	Comma-separated list of datasets for text encoder evaluation
--template	str	None	Prompt template to use for zero-shot classification

Method Descriptions

- **finetune**: Standard fine-tuning of the model
- **lwf**: Learning without Forgetting - uses distillation from previous model
- **ZSCL**: Zero-Shot Continual Learning - regularizes with reference data/text
- **icarl**: Incremental Classifier and Representation Learning - uses exemplar memory

Train Mode Descriptions

- **whole**: Train the entire model (both encoders)
- **text**: Train only the text encoder
- **image**: Train only the image encoder
- **image-fc**: Train image encoder and classification head
- **image-fc-fixed**: Train image encoder with fixed classification head
- **fc**: Train only the classification head (linear probe)

Single Image Evaluation

Argument	Type	Default	Description
--eval-single	str	None	Path to a single image for evaluation
--prompt	str	None	Custom prompt for single image evaluation
--class-names	str	None	Path to file containing class names (one per line)

Example:

```
python -m src.main \
    --load ckpt/model.pth \
    --eval-single path/to/image.jpg \
    --class-names data/text_classes/imagenet_classes.txt \
    --eval-only
```

Save & Load

Argument	Type	Default	Description
----------	------	---------	-------------

Argument	Type	Default	Description
--save	str	None	Directory path to save checkpoints
--load	str	None	Path to checkpoint file to load
--load_federate	str	None	Comma-separated paths for federated model loading

Model Control (image-fc branch)

Argument	Type	Default	Description
--fair	flag	False	Enable fair comparison mode
--we	flag	False	Enable weight ensembling
--we_wise	flag	False	Enable WiSE weight ensembling
--we_wise_alpha	float	0.98	Alpha for WiSE weight ensembling
--moving_avg	flag	False	Use exponential moving average of weights
--avg_freq	int	100	Frequency of moving average updates
--mv_avg_decay	float	0.999	Decay rate for moving average
--mv_avg_model	str	n	Base model for moving average. Choices: n, t, zeroshot
--l2	float	0	L2 regularization strength toward reference model
--fc-init	flag	False	Reinitialize the classification head
--fc-setnone	flag	False	Set classification head to None
--dataset-shift	flag	False	Enable dataset shift mode
--n_class	int	10	Number of classes for classification head

ZSCL (Zero-Shot Continual Learning)

Argument	Type	Default	Description
--ref-dataset	str	None	Reference dataset for regularization (e.g., ImageNet)
--ref-sentences	str	None	Reference text embeddings (e.g., conceptual_captions)
--ref-model	str	None	Path to reference model checkpoint
--ref-wise	flag	False	Use WiSE-FT for reference model
--ref_wise_alpha	float	0.8	Alpha for reference WiSE-FT
--T	float	2.0	Temperature for distillation loss
--num	float	64	Number of reference samples per batch

Example (ZSCL training):

```
python -m src.main \
--method ZSCL \
--train-dataset DTD \
--ref-dataset ImageNet \
--ref-sentences conceptual_captions \
--T 2.0 \
--num 64
```

iCaRL

Argument	Type	Default	Description
--dataset_order	str	None	Comma-separated list of datasets defining task order
--memory_size	int	10000	Total number of exemplars to store in memory

Example:

```
python -m src.main \
--method icarl \
--dataset_order "CIFAR10,CIFAR100/DTD" \
--memory_size 2000
```

Loss Functions

Argument	Type	Default	Description
--weight_adjust	flag	False	Enable weight adjustment for loss
--feature_mse	flag	False	Add MSE loss on features
--image_loss	flag	False	Enable image encoder loss
--text_loss	flag	False	Enable text encoder loss
--ablation_loss_2	flag	False	Enable second ablation loss variant

WiSE-FT (Weight-Space Ensembling)

Argument	Type	Default	Description
--wise_merge	flag	False	Use WiSE merging during training
--wise_ft	flag	False	Use WiSE-FT during evaluation
--wise_ft_model	str	n	Model to ensemble with. Choices: n, zeroshot

Argument	Type	Default	Description
--wise_ft_alpha	float	0.8	Interpolation factor (0=finetuned, 1=reference)
--wise-ft	flag	False	Alternative flag for WiSE-FT
--alpha	float	0.5	Alpha for --wise-ft flag

WiSE-FT Formula:

```
final_weights = alpha * finetuned_weights + (1 - alpha) * zeroshot_weights
```

Experiment Organization

Argument	Type	Default	Description
--exp_name	str	None	Name of the experiment for organization
--results-db	str	results.jsonl	Path to JSONL file for storing results
--cache-dir	str	None	Directory for caching features and encoders

Model Freezing

Argument	Type	Default	Description
--freeze-encoder	flag	False	Freeze the image encoder (for linear probing)
--freeze-fc	int	0	Number of FC layers to freeze

Learning Without Forgetting (LwF)

Argument	Type	Default	Description
--lwf	flag	False	Enable LwF distillation loss
--basic_model_load	str	None	Comma-separated paths to load base classifiers
--fc_load	str	None	Comma-separated paths to load FC layers
--keep_old_heads	int	0	Number of old classification heads to retain

Baseline & TRIO

Argument	Type	Default	Description
--baseline	flag	False	Run baseline experiment
--trio	flag	False	Enable TRIO method

Argument	Type	Default	Description
--control-dataset	str	None	Control dataset for TRIO
--control-dataset-add	str	None	Additional control dataset
noise	positional	-	Use random noise for regularization
--rff	flag	False	Enable random Fourier features

Fisher Weighting

Argument	Type	Default	Description
--fisher	str	None	Comma-separated paths to Fisher information matrices
--fisher_floor	float	1e-8	Minimum value for Fisher weights (numerical stability)

Thesis-Specific Arguments

Argument	Type	Default	Description
--freeze	flag	False	Freeze model parameters
--mixup	int	None	Enable mixup augmentation with specified alpha
--orthogonal-gradients	int	None	Number of orthogonal gradient projections
--orthogonal-gradients-path	str	None	Path(s) to orthogonal gradient basis (multiple allowed)
--untrained	flag	False	Use untrained (random) model weights
--custom-finetune	flag	False	Enable custom fine-tuning procedure
--max-evaluation-size	int	None	Limit evaluation dataset size

LoRA (Low-Rank Adaptation)

LoRA enables parameter-efficient fine-tuning by freezing the base CLIP model and training small low-rank adapter layers. This significantly reduces memory usage and training time while maintaining good performance.

Argument	Type	Default	Description
--lora	flag	False	Enable LoRA training. Freezes base model and trains only LoRA adapter layers.
--lora-r	int	8	LoRA rank (dimension of low-rank matrices). Higher = more capacity but more parameters.
--lora-alpha	int	16	LoRA scaling factor. Effective scaling is α/r .

Argument	Type	Default	Description
--lora-dropout	Float	0.1	Dropout probability applied to LoRA layers.
--lora-target-modules	str	None	Comma-separated module names to apply LoRA. Default: attention and MLP layers. ("q_proj", "k_proj", "v_proj", "out_proj")
--lora-bias	str	none	Which biases to train. Choices: none, all, lora_only

LoRA Concepts

- **Rank (r)**: Controls the capacity of LoRA adapters. Typical values: 4, 8, 16, 32. Lower rank = fewer parameters but potentially less expressive.
- **Alpha**: Scaling factor for LoRA updates. The effective learning rate scaling is alpha/r.
- **Target Modules**: By default, LoRA is applied to attention projections and MLP layers in both visual and text encoders.

Requirements

LoRA requires the `peft` library:

```
pip install peft
```

Common Usage Examples

Basic Fine-tuning

```
python -m src.main \
--method finetune \
--train-mode whole \
--train-dataset DTD \
--eval-datasets DTD,ImageNet \
--iterations 1000 \
--lr 1e-5 \
--batch-size 32 \
--save ckpt/dtd_finetune
```

ZSCL with Reference Data

```
python -m src.main \
--method ZSCL \
--train-mode whole \
--train-dataset DTD \
--ref-dataset ImageNet \
```

```
--ref-sentences conceptual_captions \
--iterations 1000 \
--lr 1e-5 \
--T 2.0 \
--save ckpt/dtd_zscl
```

Evaluation Only

```
python -m src.main \
--load ckpt/model.pth \
--eval-datasets DTD,CIFAR100,ImageNet \
--eval-only
```

WiSE-FT Evaluation

```
python -m src.main \
--load ckpt/model.pth \
--eval-datasets ImageNet \
--eval-only \
--wise_ft \
--wise_ft_alpha 0.5
```

Linear Probe (FC only)

```
python -m src.main \
--method finetune \
--train-mode fc \
--freeze-encoder \
--train-dataset CIFAR100 \
--n_class 100 \
--iterations 5000 \
--save ckpt/cifar100_probe
```

LoRA Fine-tuning

```
# Basic LoRA training (default settings)
python -m src.main \
--method finetune \
--train-mode whole \
--lora \
--train-dataset DTD \
--iterations 1000 \
--lr 1e-4 \
--save ckpt/dtd_lora
```

```
# LoRA with custom configuration
python -m src.main \
    --method finetune \
    --train-mode whole \
    --lora \
    --lora-r 16 \
    --lora-alpha 32 \
    --lora-dropout 0.05 \
    --train-dataset DTD \
    --iterations 1000 \
    --lr 1e-4 \
    --save ckpt/dtd_lora_r16

# LoRA with ZSCL (continual learning)
python -m src.main \
    --method ZSCL \
    --train-mode whole \
    --lora \
    --lora-r 8 \
    --train-dataset DTD \
    --ref-dataset ImageNet \
    --iterations 1000 \
    --save ckpt/dtd_lora_zscl
```

Notes

- `--epochs` and `--iterations` are mutually exclusive; use one or the other
- `--eval-interval` and `--eval-every-epoch` are mutually exclusive
- Device is automatically set to CUDA if available, otherwise CPU
- For ZSCL, both `--ref-dataset` and `--ref-sentences` are typically used together for best results